



Integration platform for devs

The easiest way for developers to connect their stack and automate workflows.

Pipedream.com

a on Ubuntu 20.04

stall Apache Kafka on Ubuntu 20.04

LTS or any Cloud Instance like Amazon EC2, Azure VM, Google Compute Engine, etc. with preinstalled Ubuntu.

Table of Contents



Introduction

Kafka is most popular open source software which provides a framework for storing, reading and analyzing streaming data. To know more about Kafka visit official [documentation page](#).

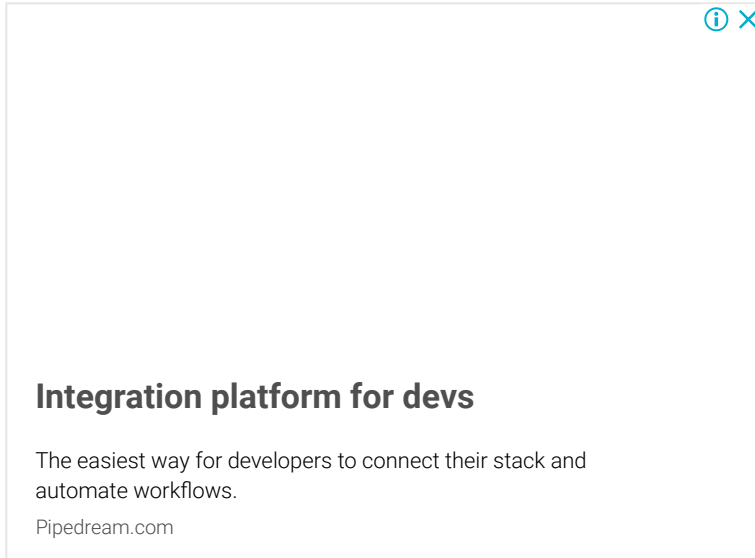
Prerequisites

- Ubuntu 20.04 with Minimal Installation
- SSH Access with Sudo privileges
- Firewall Ports: 2181, 9092
- JDK 1.8 or higher version



Install JDK on Ubuntu

How to Install Oracle JAVA 8 on Ubuntu 20.04



```
$ sudo apt install openjdk-8-jdk
```

```
$ sudo apt install openjdk-11-jdk
```

To check the java version, Here i have install installed Oracle Java 8 on my system.

```
$ java -version
```

Output:

```
java version "1.8.0_251"
```


```
Java(TM) SE Runtime Environment (build 1.8.0_251-b07)
```

```
Java HotSpot(TM) 64-Bit Server VM (build 25.251-b07, mixed mode)
```

Step 1: How to Install Apache Kafka on Ubuntu 20.04 LTS

To download the binary from official website. Please use this [link to download](#) and to prompts to download page.





Integration platform for devs

The easiest way for developers to connect their stack and automate workflows.

Pipedream.com

Join our community of 50k+ developers processing millions of events

Pipedream.com

kafka/2.4.1/kafka_2.13-2.4.1.tgz

Now to un-tar the archive file and move to another location:

```
$ sudo tar xzf kafka_2.13-2.4.1.tgz
```

```
$ sudo mv kafka_2.13-2.4.1 /opt/kafka
```

Step 2: Creating zookeeper and Kafka Systemd Unit Files

Create the systemd unit file for zookeeper service

```
$ sudo nano /etc/systemd/system/zookeeper.service
```

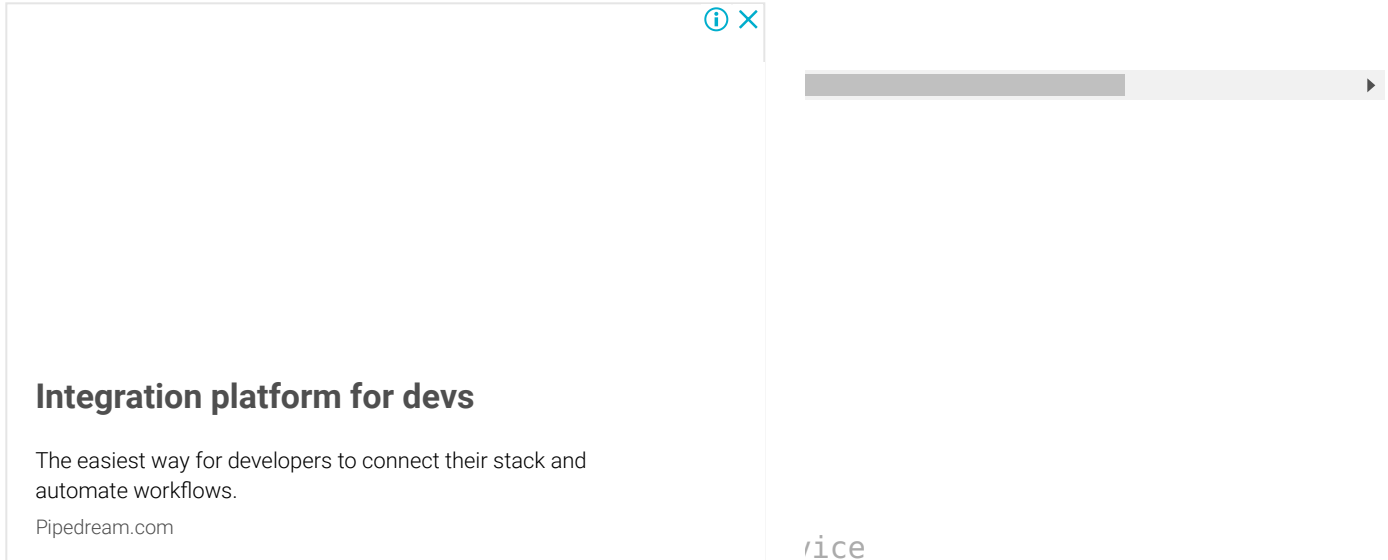
paste the below lines

```
/etc/systemd/system/zookeeper.service
```

```
[Unit]
Description=Apache Zookeeper service
Documentation=http://zookeeper.apache.org
Requires=network.target remote-fs.target
After=network.target remote-fs.target

[Service]
Type=simple
ExecStart=/opt/kafka/bin/zookeeper-server-start.sh /opt/kafka/conf/zooe
```

```
ExecStop=/opt/kafka/bin/zookeeper-server-stop.sh
Restart=on-abnormal
```



paste the below lines

```
/etc/systemd/system/kafka.service
```

```
[Unit]
Description=Apache Kafka Service
Documentation=http://kafka.apache.org/documentation.html
Requires=zookeeper.service

[Service]
Type=simple
Environment="JAVA_HOME=/opt/jdk/jdk1.8.0_251"
ExecStart=/opt/kafka/bin/kafka-server-start.sh /opt/kafka/config/server.pr
ExecStop=/opt/kafka/bin/kafka-server-stop.sh

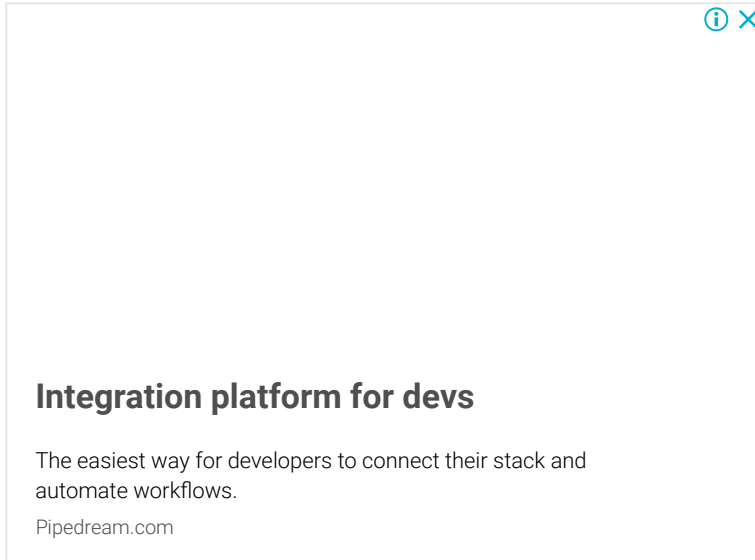
[Install]
WantedBy=multi-user.target
```

Note: Modify the value of **JAVA_HOME** value, If the path of your Java installation is different.

Reload the daemon to take effect

```
$ sudo systemctl daemon-reload
```

Step 3: Start ZooKeeper and Kafka Service



```
$ sudo systemctl status zookeeper
```

Output:

```
● zookeeper.service - Apache Zookeeper service
   Loaded: loaded (/etc/systemd/system/zookeeper.service; disabled; vendor pr
   Active: active (running) since Sat 2020-06-11 09:52:40 UTC; 22min ago
   Docs: http://zookeeper.apache.org
   Main PID: 25638 (java)
   Tasks: 30 (limit: 9513)
   CGroup: /system.slice/zookeeper.service
           └─25638 java -Xmx512M -Xms512M -server -XX:+UseG1GC -XX:MaxGCPauseMillis=2
```


Check the status of kafka service if it started

```
$ sudo systemctl status kafka
```


Output:

```
● kafka.service - Apache Kafka Service
   Loaded: loaded (/etc/systemd/system/kafka.service; disabled; vendor preset
   Active: active (running) since Sat 2020-06-11 09:52:48 UTC; 24min ago
   Docs: http://kafka.apache.org/documentation.html
   Main PID: 26027 (java)
   Tasks: 66 (limit: 9513)
   CGroup: /system.slice/kafka.service
```

```
CGroup: /system.slice/kafka.service
└─26027 /opt/jdk/jdk1.8.0_241/bin/java -Xmx1G -Xms1G -server -XX:+UseG1GC
```



**server in the
stemd unit file)**



Integration platform for devs

The easiest way for developers to connect their stack and automate workflows.
Pipedream.com



in our community of 50k+
developers processing
ons of events
edream.com

We have shell script to run the kafka and zookeeper server in backend:

Create a file named kafkastart.sh and copy the below script:

```
#!/bin/bash
sudo nohup /opt/kafka/bin/zookeeper-server-start.sh -daemon /opt/kafka/con
sleep 5
sudo nohup /opt/kafka/bin/kafka-server-start.sh -daemon /opt/kafka/config/
```

After give the executable permissions to the file:

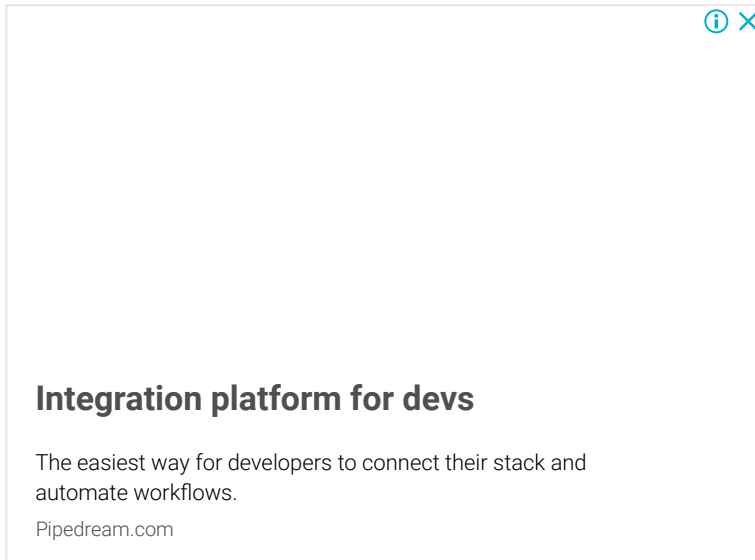
```
sudo chmod +x kafkastart.sh
```

Successfully We have covered how to install apache kafka on ubuntu 20.04 LTS.

Step 4: Creating Topic in Kafka



Now we will create a topic named as “DevOps” with a single replicaton-factor and partition.



```
keeper localhost:2181 --replicati
```

To check the list of topics created.

```
$ sudo bin/kafka-topics.sh --list --zookeeper localhost:2181
```

Output:

```
DevOps
```

Step 5: To send some messages

To send some messages for created Topic.

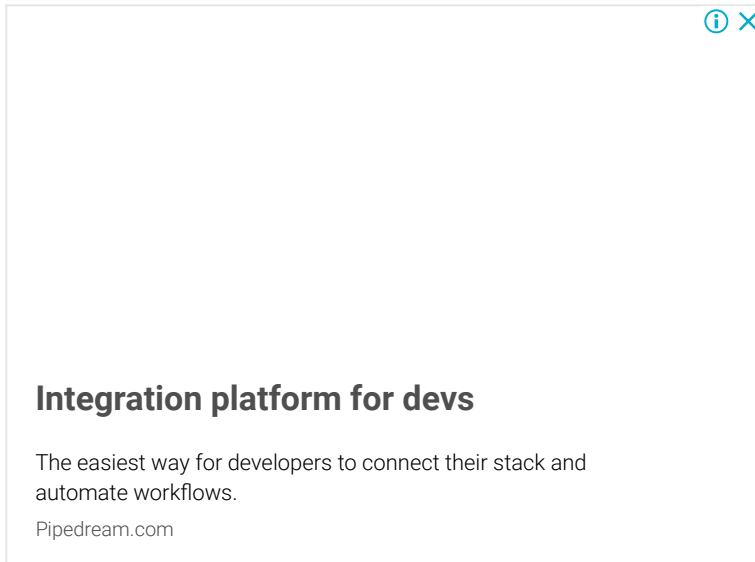
```
$ sudo bin/kafka-console-producer.sh --broker-list localhost:9092 --topic
```

its prompt for messages to type:

```
> Hi  
> How are you?
```

Step 6: To Start consumer

Using below command we can see the list of messages:



```
bootstrap-server localhost:9092 --t
```

remote machine

on remote server. Please follow

```
cd /opt/kafka/config
```

Now look for server.properties and make some configuration changes:

```
sudo vi server.properties
```

In this properties file uncomment as mentioned below:

```
listeners=PLAINTEXT://:9092
```

```
advertised.listeners=PLAINTEXT://<HOST IP>:9092
```

Step 8: To Delete any Topic

If you want to delete any created topic use below command:

```
$ sudo bin/kafka-topics.sh --delete localhost:9092 --topic <anytopic>
```

Conclusion:

In this article, We have performed ,How to Install Apache Kafka on Ubuntu 20.04 LTS and also covered creating customd unit file for zookeeper and kafka service, start

c, start kafka and zookeeper

Integration platform for devs

The easiest way for developers to connect their stack and automate workflows.

Pipedream.com

.04 LTS

TS



Sivasai Sagar

I am working as DevOps Engineer and having 5 years of Experience.
Likes to share knowledge.

Leave a Comment





Integration platform for devs

The easiest way for developers to connect their stack and automate workflows.

Pipedream.com

w your comment data is processed.



Recent Posts

[Shell Script to Build Docker Image \[2 Steps\]](#)

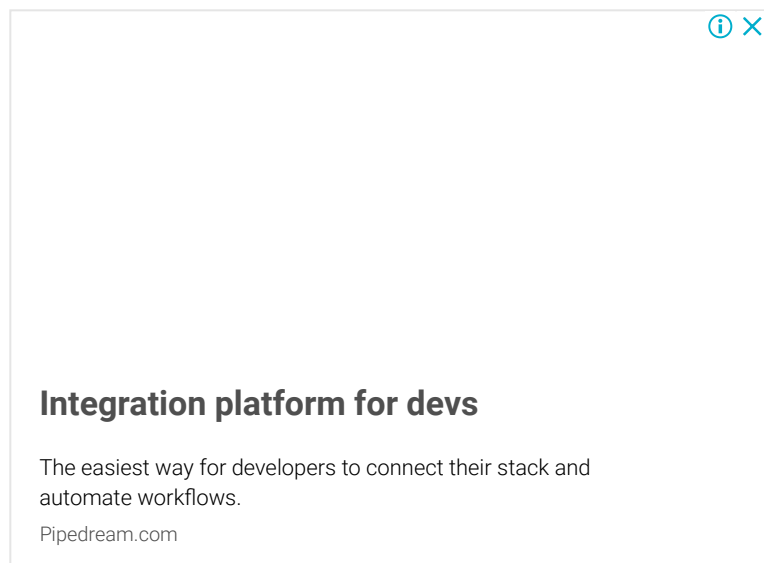
[4 Steps to Install Kubernetes Dashboard](#)

[Shell Script to Start MongoDB Service if Down \[2 Steps\]](#)

[9 Steps to Setup Kubernetes on AWS using KOPS](#)

[How to Install SonarQube on Ubuntu 20.04 LTS](#)





About FOSS TechNix

FOSS TechNix (Free ,Open Source Softwares and Technology Nix*) is a community site where you can find How-to Guides, Articles,Tips and Tricks for DevOps Tools,Linux ,Databases,Clouds and Automation.

Site Links

[About Us](#)

[Privacy Policy](#)

[Contact Us](#)

[Terms and Conditions](#)

Archives



Select Month



cookies to improve your experience. By using this
Privacy Policy

Integration platform for devs

The easiest way for developers to connect their stack and
automate workflows.

Pipedream.com

